


# R functional programming

---

Alberto Garfagnini

Università di Padova

Advanced R 02



## Function fundamentals

---

- R functions can be broken into 3 components:
  - **arguments** : the list of arguments that describe how to call the function
  - **body** : the code inside the function
  - **environment** : the data structure that tell us how the function finds the values associated with the name

```
mysum <- function(x, y) {  
  # Compute the sum of 2 vectors  
  x + y  
}
```

```
> formals(mysum)          body(mysum)  
#> $x                     #> {  
#> $y                     #>   x + y  
                          #> }
```

```
environment(mysum)  
#> <environment: R_GlobalEnv>
```

- functions, as objects, can have attributes

```
attributes(mysum)  
#> $srcref  
#> function(x, y) {  
#>   # Compute the sum of 2 vectors  
#>   x + y  
#> }  
  
attr(mysum, "srcref")  
#> function(x, y) {  
#>   # Compute the sum of 2 vectors  
#>   x + y  
#> }
```

- are those found in the base package
- are primarily written in C, so their `formals()`, `body` and `environment()` are all `NULL`

```
sum
#> function (... , na.rm = FALSE) .Primitive("sum")

formals(sum)
#> NULL

body(sum)
#> NULL

environment(sum)
#> NULL

typeof(sum)
#> [1] "builtin"
```

## Creating functions

---

### A "named" function

- 1) create a function object with `function`
- 2) bind it to a name with `<-`

```
mym <- function(x) {
  sin(1 / x ^ 2)
}
mym(1:4)
#> [1] 0.84147098 0.24740396 0.11088263 0.06245932
```

### Anonymous functions

- it is done when a function name (i.e. binding) is not given

```
integrate(function(x) sin(x) ^ 2, 0, pi)
#> 1.570796 with absolute error < 1.7e-14
```

### List of functions

- functions can be put in a list

```
lfuns <- list(
  half = function(x) x/2,
  double = function(x) x*2)
lfuns$half(10)
#> [1] 5
lfuns$double(10)
#> [1] 20
```

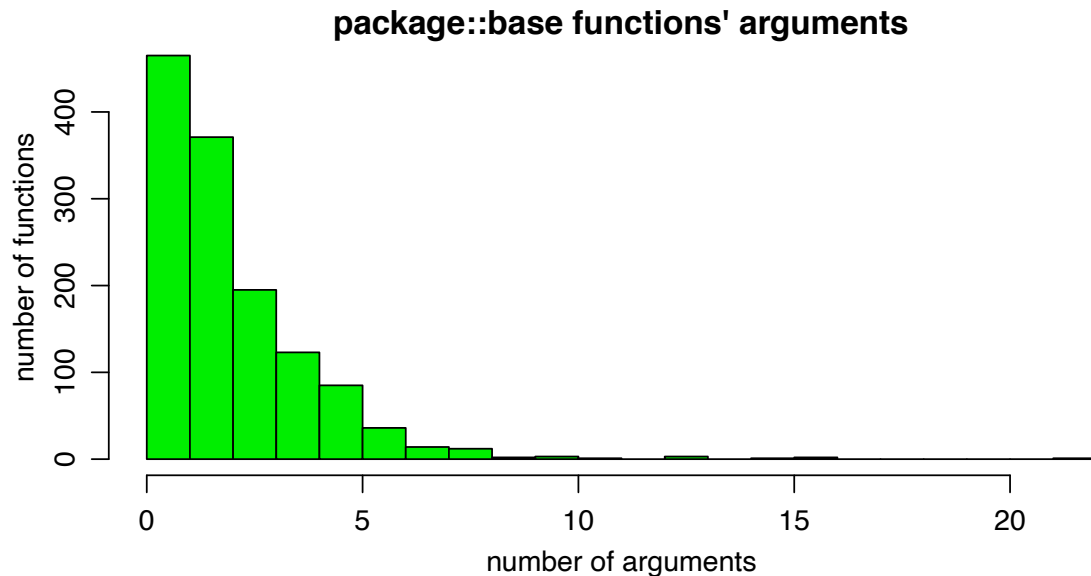
# Exercise

- the following code create a list of all functions in the base package

```
objs <- mget(ls("package:base", all=TRUE), inherits=TRUE)
bfuns <- Filter(is.function, objs)
```

1→ Determine the number of arguments for all functions and plot the distributions

2→ How to restrict the search only to primitive functions ?



A. Garfagnini (UniPD)

AdvStat 4 PhysAna - RAdv01

4

## Functions calling

- R functions are normally invoked by placing the arguments in parentheses:

```
x <- c(1:3, NA, 5:10)
mean(x, na.rm=TRUE)
#> [1] 5.666667
```

- in case the functions arguments are inside a data structure
- the `do.call()` function can be called, instead:

```
x <- c(1:3, NA, 5:10)
args <- list(x, na.rm=TRUE)
do.call(mean, args)
#> [1] 5.666667
```

## Functions composition

- let's imagine we need to call several functions:

```
square <- function(x) x^2
deviation <- function(x) x - mean(x)
x <- runif(10^3)
```

- we can nest the function calls

```
sqrt(mean(square(deviation(x))))
#> [1] 0.2925719
```

A. Garfagnini (UniPD)

AdvStat 4 PhysAna - RAdv01

5

## Functions calling (2)

- we could also store intermediate results as vectors

```
out <- deviation(x)
out <- square(out)
out <- mean(out)
out <- sqrt(out)
out
#> [1] 0.2925719
```

- but we could also use the pipe operator, %>%

```
library(magrittr)

x %>%
  deviation() |>
  square() |>
  mean() |>
  sqrt()
#> [1] 0.2925719
```

- $x \mid > f()$  is equivalent to  $f(x)$
- $x \mid > f(y)$  is equivalent to  $f(x, y)$

## Lazy evaluation

---

- all function arguments are **lazy evaluated**

```
hstop <- function(x) { 10 }

hstop(1)
#> [1] 10

hstop(stop("This is an error!"))
#> [1] 10

stop("This is an error!")
#> Error: This is an error!
```

## Promises

- unevaluated argument is called a promise, or a thunk.
- a promise is made up of two parts:
  - an expression, line  $x + y$  which gives rise to delayed computation
  - an environment, where the expression should be evaluated

# Function arguments : default values

---

- function arguments can have default values

```
f <- function(a = 1, b = 2) c(a, b)
f()
#> [1] 1 2
```

- since arguments are evaluated lazily, default arguments can be defined in terms of other arguments

```
g <- function(a = 1, b = a * 2) c(a, b)
g()
#> [1] 1 2
g(10)
#> [1] 10 20
```

- if an argument was supplied or not can be seen with the `missing()` function

```
i <- function(a, b) { c(missing(a), missing(b)) }
i()
#> [1] TRUE TRUE
i(a=1)
#> [1] FALSE TRUE
i(b=1)
#> [1] TRUE FALSE
i(1,2)
#> [1] FALSE FALSE
```

## The ... (dot-dot-dot) function argument

---

- it is a special argument called ...
- it will match any arguments not otherwise matched, and can be easily passed on to other functions
- one relatively sophisticated user of ... is the base `plot()` function
- `plot()` is a generic method with arguments `x`, `y` and ...
- simple invocations of `plot()` end up calling `plot.default()` which has many more arguments (including ...). In this way, `plot()` accepts graphical parameters which are listed in the help of `par()`

```
plot(1:5, col = "red")
plot(1:5, cex = 5, pch = 20)

# The following allows to capture the arguments
f <- function(...) {
  names(list(...))
}
f(alpha=1, slope=3)
[1] "alpha" "slope"
```

# Every operation is a function call

---

## Golden rules

- everything that exists in R is an object
- but everything that happens is a function call
- this includes infix operators like `+`, control flow operators like `for`, `if`, and `while`, subsetting operators like `[]` and `$`, and even the curly brace `{`
- the backtick lets us refer to functions or variables that have otherwise reserved or illegal names

```
x <- 10; y <- 5; x + y  
[1] 15
```

```
`+`(x, y)  
[1] 15
```

```
for (i in 1:2) print(i)  
[1] 1  
[1] 2
```

```
`for`(i, 1:2, print(i))  
[1] 1  
[1] 2
```

```
> { print(1) }  
[1] 1  
> `{(print(1))  
[1] 1
```

# Every operation is a function call

---

- this allows to override the definitions of these special functions
- usually it is a bad idea, but it allows you to do something that would have otherwise been impossible
- example: we need to add 3 to every element of a list
- option 1: define a function `add()` and use `sapply()`:

```
add <- function(x, y) x + y  
sapply(1:10, add, 3)  
[1] 4 5 6 7 8 9 10 11 12 13
```

- but we can also get the same effect using the built-in `+` function:

```
sapply(1:5, `+`, 3)  
[1] 4 5 6 7 8
```

```
sapply(1:5, "+", 3)  
[1] 4 5 6 7 8
```

- the second version works as well, because `sapply()` can be given the name of a function instead of the function itself
- it uses `match.fun()` to find functions given their names

# Function arguments

---

- it is useful to distinguish between
- **formal arguments** → a property of the function
- **actual arguments** → can vary each time you call the function
- when calling a function, arguments can be specified by
  - position, complete name, partial name
- arguments are matched first by exact name (perfect matching), then by prefix matching, and finally by position

```
f <- function(alpha, beta1, beta2) {  
  list(a = alpha, b1 = beta1, b2 = beta2)  
}  
str(f(1,2,3))  
List of 3  
 $ a : num 1    $ b1: num 2    $ b2: num 3  
  
str(f(2,3,alpha=1))  
List of 3  
 $ a : num 1    $ b1: num 2    $ b2: num 3  
  
str(f(2,3,a=1))  
List of 3  
 $ a : num 1    $ b1: num 2    $ b2: num 3  
  
str(f(1,2,beta=3))  
Error in f(1, 2, beta = 3) : argument 3 matches multiple formal arguments
```

## Special calls: Infix functions

---

- most functions in R are *prefix* operators: the name of the function comes before the arguments
- infix functions are those where the function name comes in between its arguments (for instance '+' or '-')
- all user created infix functions must start and end with %
- R comes with the following infix functions predefined: %, %\*, %/, %in%, %o%, %x%
- the complete list of built-in infix operators that don't need % is: ::, :::, \$, , ^, \*, /, +, -, >, >=, <, <=, ==, !=, !, &, &&, |, ||, ~, <-, <<-
- we could create a new operator that pastes together strings:

```
`%+%` <- function(a, b) paste(a, b, sep = "")  
"new" +% "string"  
[1] "newstring"
```

- as far as R is concerned there is no difference between these two expressions:

```
"new" +% "string"  
[1] "newstring"  
`%+%`("new", "string")  
[1] "newstring"
```

# Special calls: replacement calls

---

- they act like they modify their arguments in place, and have the special name `xxx <-`
- they typically have two arguments (x and value), although they can have more, and they must return the modified object

```
`second<-` <- function(x, value) {  
  x[2] <- value  
  x  
}  
x <- 1:5  
second(x) <- 0  
x  
[1] 1 0 3 4 5
```

- when R evaluates the assignment `second(x) <- 5`, it notices that the left hand side of `<-` is not a simple name, so it looks for a function named `second<-` to do the replacement
- if additional arguments are needed, they go in between x and value

```
`modify<-` <- function(x, position, value) {  
  x[position] <- value  
  x  
}  
modify(x, 1) <- -5  
x  
[1] -5 0 3 4 5
```

## Functions : additional topics

---

### Return values

- the last expression evaluated in a function becomes the return value

```
f <- function(x) {  
  if ( x < 10 ){ 0 } else { 10 }  
}  
f(5)  
[1] 0
```

- functions can return only a single object
- this is not a limitation because they can return a list containing any number of objects

### Invisible values

- functions can return invisible values, which are not printed out by default when you call the function

```
f1 <- function() 1  
f2 <- function() invisible(1)  
f1()  
[1] 1  
f2()  
f1() == 1  
[1] TRUE  
f2() == 1  
[1] TRUE
```

- the most common function that returns invisibly is `<-`



# Functions: `on.exit()` trigger

---

- functions can set up other triggers to occur when the function is finished using `on.exit()`
- the code inside `on.exit()` is always run, regardless of how the function exits, whether with an explicit (early) return, an error, or simply reaching the end of the function body

```
in_dir <- function(dir, code) {  
  old <- setwd(dir)  
  on.exit(setwd(old))  
  force(code)  
}  
  
getwd()  
[1] "/Users/alberto/Documents/didattica/PhysicsOfData/R_code"  
  
in_dir("~", getwd())  
[1] "/Users/alberto"  
  
getwd()  
[1] "/Users/alberto/Documents/didattica/PhysicsOfData/R_code"
```

## Functionals basics

---

### Definition

a **FUNCTIONAL** is a function that  
takes **FUNCTION as INPUT**  
and returns a **VECTOR as OUTPUT**

- example:

```
randomize <- function(f) f(runif(10^3))  
  
randomize(mean)  
#> [1] 0.4954407  
randomize(mean)  
#> [1] 0.491658  
  
randomize(sum)  
#> [1] 507.5148
```

- typical examples in base R:  
`lapply()`, `apply()` and `tapply()`
- other example: `integrate()`

```
integrate(dnorm, -Inf, Inf)  
#> 1 with absolute error < 9.4e-05
```

# Functionals : replacement for loops ?


a common use of **functionals** is as **alternative to for loops**

## NOTE

- for loops are not slow by themselves
- what makes them slow is **what programmers do inside the for loop body**

ex: modifying a data structure makes the loop slow because each modification creates a copy: copy-on-modify

functionals  
for  
while  
repeat



- switching from loop to functional is a pattern matching exercise:  
goal: **find a functional that matches the basic loop form**

## Our first functions: **purrr::map()**

- it takes a vector and a function
- it calls the function for each vector element
- it returns the results in a list


```
purrr::map(1:10, f)
```

is equivalent to

```
list(f(1), f(2), ..., f(10))
```

```
double <- function(x) x*2
xd <- purrr::map(1:10, double)
str(xd)
#> List of 10
#> $ : num 2
#> $ : num 4
...
#> $ : num 18
#> $ : num 20

unlist(xd)
#> [1] 2 4 6 8 10 12 14 16 18 20
```



# Example 1

- we have a tibble with different data sets

```
dt <- tibble( a1 = rnorm(10), b1 = runif(10),  
              c1 = rpois(10, 3.7), d1 = rbeta(10, 0.3, 5) )
```

- we want to evaluate the median of each column

```
omed <- vector("double", ncol(dt))  
omed  
#> [1] 0 0 0 0  
for (i in seq_along(dt)) {  
  omed[[i]] <- median(dt[[i]])  
}  
omed  
#> [1] 0.165312063 0.487255521 4.000000000 0.009203981
```

- it's possible to wrap up for loops in a function, and call that function instead of using the for loop directly

```
purrr::map_dbl(dt, median)  
#>      a1      b1      c1      d1  
#> 0.165312063 0.487255521 4.000000000 0.009203981
```

- all the map\_\*() functions use ... to pass along additional arguments to .f each time it's called


```
purrr::map_dbl(dt, mean, trim=0.5)  
#>      a1      b1      c1      d1  
#> 0.165312063 0.487255521 4.000000000 0.009203981
```

## map()

the function `map()` returns a list:

```
      a list or vector  
      ||  
      \/  
map(.x, .f, ...)  
  /\  |-----> pass additional arguments to .f each  
  ||      time it is called  
  ||  
      a function, formula or vector
```

- `map_lgl()`, `map_int()`, `map_dbl()` and `map_chr()` return a vector of specific type (logical, integer, double or character)
- `map_dfr()` and `map_dfc()` return a data frame created by row or by column
- any arguments that come after `f` in the call to `map()` are inserted after the data in individual calls to `f()`



## Example 2: `map()`

- we generate 10 sets of random numbers from a probability distribution

```
1:10 %>% map(rnorm, n=20) -> 11
```

- this can be done using an anonymous function

```
1:10 %>% map(function(x) rnorm(n=20, x)) -> 12
```

- or by using a **one-sided formula**

```
1:10 %>% map( ~ rnorm(n=20, .x) ) -> 13
```

- there are a few shortcuts that you can use with `.f` in order to save a little typing
  - `.x` and `.y` are used for two argument functions, and `..1`, `..2`, `..3`, ... for all the additional arguments
- `map()` can be chained:

```
1:10 %>%  
  map(rnorm, n=20) %>%  
  map_dbl(mean)  
#> [1] 0.8355395 2.0266397 3.1451209 3.9854774 5.0977312  
#> [6] 6.0904780 6.9547342 8.4906865 8.9917292 10.1268192
```

## Mapping over multiple arguments: `map2()`


- as an example we want to generate several data sets from a normal distribution with different mean and variance

```
mus <- list(5, 10, -3)  
sigmas <- list(1, 5, 10)  
map2(mus, sigmas, rnorm, n = 5) %>% str()  
#> List of 3  
#> $ : num [1:5] 4.17 5.24 5.54 4.8 5.44  
#> $ : num [1:5] 12.71 7.01 9.56 7.25 10.74  
#> $ : num [1:5] -8.72 9.89 -14.54 3.51 -9.49
```

- the same results could have been done iterating over indices

```
seq_along(mus) %>%  
  map( ~ rnorm(5, mus[[.]], sigmas[[.]] ) ) %>% str()  
#> List of 3  
#> $ : num [1:5] 4.73 6.52 2.68 5.42 5.35  
#> $ : num [1:5] 14.913 -0.695 11.702 17.911 1.795  
#> $ : num [1:5] -4.14 -1.08 -7.85 5.79 13.73
```

- but the **code with `map2()` is simpler and cleaner**




## additional functions: `pmap()` and `imap()`

- in case of multiple arguments, `purrr` provides `pmap()` which takes a list of arguments
- if you don't name the list's elements, `pmap()` will use positional matching when calling the function. This makes the code harder to read → use named arguments:

```
args2 <- list(mean = c(5, 10, -3),
              sd = c(1, 5, 10), n = 5)
args2 %>%
  pmap(rnorm) %>%
  str()

#> List of 3
#> $ : num [1:5] 5.38 4.54 3.85 5.44 5.42
#> $ : num [1:5] 17.038 6.072 15.107 0.697 6.488
#> $ : num [1:5] -10.7 2.99 -1.12 17.47 13.08
```



## Invoking different functions: `invoke_map()`

- a setup in complexity is to invoke different functions with different parameters (values and meanings):

```
fgen <- c("runif", "rnorm", "rpois")

fpar <- list(
  list(min = -1, max = 1),
  list(sd = 3),
  list(lambda = 7.5))

invoke_map(fgen, fpar, n = 5) %>% str()

#> List of 3
#> $ : num [1:5] -0.7744 -0.0524 0.7523 0.5074 0.5284
#> $ : num [1:5] 3.162 -2.766 -0.298 -2.849 -2.638
#> $ : int [1:5] 5 7 10 6 4
```

- our data is organized in text files accoring to different years:
  - data\_2020\_Italy.csv, data\_2021\_Italy.csv
- we want to read the data and combine them in one data.frame

```
read_my_csv <- function(year, country) {  
  filename <- paste0(year, "_", country, ".csv")  
  mobdata_dir <- "./Region_Mobility_Report_CSVs"  
  filepath <- file.path(mobdata_dir, filename)  
  message(paste("Reading from file:", filepath))  
  read_csv(filepath)  
}  
  
years <- 2020:2021  
country <- "Italy"  
  
mbdata <- map_df(years, read_my_csv, country)  
  
Reading from file: ./Region_Mobility_Report_CSVs/2020_IT.csv  
...  
Reading from file: ./Region_Mobility_Report_CSVs/2021_IT.csv  
...
```