

EslamiShafigh_Ehsan_rlab05

Ehsan Eslmai Shafigh

2023-05-26

Assignment 5

Exercise 1

We calculate a different posterior for each of the two sets of observations.

(a.1) Assuming a uniform prior:

```
sigma_y = 65 * 1 + 22 * 2 + 3 * 3 + 1 * 4
n = 109 + 65 + 22 + 3 + 1
n.sample = 1000
delta.lambda = 1/n.sample

alpha = 1 + sigma_y
mu = n
lambda = seq(from = 0 , by = delta.lambda , length.out = n.sample )

lambda_posterior = dgamma(lambda , alpha , mu)

mean = delta.lambda * sum(lambda * lambda_posterior)
variance = delta.lambda * sum(lambda ^ 2 * lambda_posterior) - mean ^ 2
median = qgamma(0.5 , alpha , mu )
ci = qgamma(c(0.025 , 0.975) , alpha , mu)

print(mean)

## [1] 0.615

print(variance)

## [1] 0.003075

print(median)

## [1] 0.6133341
```

```
print(ci)
```

```
## [1] 0.5111251 0.7283408
```

```
# Plot the lambda and lambda_posterior
```


```
plot(lambda, lambda_posterior, type = 'l', col = 'blue', xlab = "Lambda", ylab = "Posterior Density", m
```

```
# Add vertical dashed lines for the confidence interval
```

```
abline(v = ci[1], lty = 'dashed', col = 'red')
```

```
abline(v = ci[2], lty = 'dashed', col = 'red')
```

Posterior Distribution of Lambda



(b.1) We repeat the same task for Jeffery's prior:

```
sigma_y = 65 * 1 + 22 * 2 + 3 * 3 + 1 * 4
```

```
n = 109 + 65 + 22 + 3 + 1
```

```
n.sample = 1000
```

```
delta.lambda = 1/n.sample
```

```
alpha = 0.5 + sigma_y
```

```
mu = n
```

```
lambda = seq(from = 0 , by = delta.lambda , length.out = n.sample )
```

```
lambda_posterior = dgamma(lambda , alpha , mu)
```

```

mean = delta.lambda * sum(lambda * lambda_posterior)
variance = delta.lambda * sum(lambda ^ 2 * lambda_posterior) - mean ^ 2
median = qgamma(0.5 , alpha , mu )
ci = qgamma(c(0.025 , 0.975) , alpha , mu)

print(mean)

```

```
## [1] 0.6125
```

```
print(variance)
```

```
## [1] 0.0030625
```

```
print(median)
```

```
## [1] 0.6108341
```

```
print(ci)
```


```
## [1] 0.5088464 0.7256196
```

```

# Plot the lambda and lambda_posterior
plot(lambda, lambda_posterior, type = 'l', col = 'blue', xlab = "Lambda", ylab = "Posterior Density", m
# Add vertical dashed lines for the confidence interval
abline(v = ci[1], lty = 'dashed', col = 'red')
abline(v = ci[2], lty = 'dashed', col = 'red')

```

Posterior Distribution of Lambda



(a.2) We can repeat the same for the second set of observations:

```
sigma_y = 91 * 1 + 32 * 2 + 11 * 3 + 2 * 4
n = 144 + 91 + 32 + 11 + 2
n.sample = 1000
delta.lambda = 1/n.sample

alpha = 1 + sigma_y
mu = n
lambda = seq(from = 0 , by = delta.lambda , length.out = n.sample )

lambda_posterior = dgamma(lambda , alpha , mu)

mean = delta.lambda * sum(lambda * lambda_posterior)
variance = delta.lambda * sum(lambda ^ 2 * lambda_posterior) - mean ^ 2
median = qgamma(0.5 , alpha , mu )
ci = qgamma(c(0.025 , 0.975) , alpha , mu)

print(mean)
```

```
## [1] 0.7035714
```

```
print(variance)
```

```
## [1] 0.002512784
```

```
print(median)
```

```
## [1] 0.7023813
```

```
print(ci)
```

```
## [1] 0.6087487 0.8051570
```


```
# Create the plot
```

```
plot(lambda, lambda_posterior, type = 'l', col = 'blue', xlab = "Lambda", ylab = "Posterior Density")
```

```
# Add vertical dashed lines for the confidence interval
```

```
abline(v = ci[1], lty = 'dashed', col = 'red')
```

```
abline(v = ci[2], lty = 'dashed', col = 'red')
```



```

sigma_y = 91 * 1 + 32 * 2 + 11 * 3 + 2 * 4
n = 144 + 91 + 32 + 11 + 2
n.sample = 1000
delta.lambda = 1/n.sample

alpha = 0.5 + sigma_y
mu = n
lambda = seq(from = 0 , by = delta.lambda , length.out = n.sample )

lambda_posterior = dgamma(lambda , alpha , mu)

mean = delta.lambda * sum(lambda * lambda_posterior)
variance = delta.lambda * sum(lambda ^ 2 * lambda_posterior) - mean ^ 2
median = qgamma(0.5 , alpha , mu )
ci = qgamma(c(0.025 , 0.975) , alpha , mu)

print(mean)

```

(b.2)

```
## [1] 0.7017857
```

```
print(variance)
```

```
## [1] 0.002506402
```

```
print(median)
```

```
## [1] 0.7005956
```

```
print(ci)
```


```
## [1] 0.6070878 0.8032465
```

```

# Create the plot
plot(lambda, lambda_posterior, type = 'l', col = 'blue', xlab = "Lambda", ylab = "Posterior Density")

# Add vertical dashed lines for the confidence interval
abline(v = ci[1], lty = 'dashed', col = 'red')
abline(v = ci[2], lty = 'dashed', col = 'red')

```



Exercise 2

I use the code written for the MCMC algorithm in one of the slides:

```
# Parameters:
# func : a function whose first argument is a real vector of parameters
# func returns a log10 of the likelihood function
# theta.init : the initial value of the Markov Chain (and of func)
# n.sample: number of required samples
# sigma : standar deviation of the gaussian MCMC sampling pdf

metropolis.1dim <- function(func , theta.init , n.sample , sigma) {
  theta.cur <- theta.init
  func.Cur <- func(theta.cur)
  func.Samp <- matrix(data=NA, nrow=n.sample , ncol=2+1)
  n.accept <- 0
  rate.accept <- 0.0

  for (n in 1:n.sample) {
    theta.prop <- rnorm(n=1, mean = theta.cur, sigma)
    func.Prop <- func(theta.prop)
    logMR <- func.Prop - func.Cur # Log10 of the Metropolis ratio
    if ( logMR >=0 || logMR >log10(runif(1)) ) {
```

```

theta.cur <- theta.prop
func.Cur <- func.Prop
n.accept <- n.accept + 1
}
func.Samp[n, 1] <- func.Cur
func.Samp[n, 2] <- theta.cur
func.Samp[n, 3] <- n
}
return(func.Samp)
}

```

Now we need to define the posterior function we are going to sample from:

```

testfunc <- function(lambda) {
return(dgamma(lambda , shape = alpha , rate = mu))}

testfunc.metropolis <- function(lambda) {
return(log10(testfunc(lambda)))}

```

Now we are ready to sample from the desired posterior:

```

lambda.init <- 0.1
sample.sig <- 0.1
n.sample <- 10^5
demo <- TRUE

set.seed(20190513)
chain <- metropolis.1dim(func=testfunc.metropolis ,
theta.init = lambda.init ,
n.sample = n.sample ,
sigma = sample.sig)

```

```

mcmc.data = chain[1000:nrow(chain) , 2]


```

```

hist(mcmc.data , breaks = 1000 , freq = F , main = 'Histogram of the Samples' , xlab = 'sample values')

```


Histogram of the Samples



```
mean = mean(mcmc.data)
variance = mean(mcmc.data ^ 2) - mean ^ 2
median = quantile(mcmc.data , 0.5)
ci = quantile(mcmc.data , c(0.025 , 0.975))

print(mean)
```

```
## [1] 0.7018181
```

```
print(variance)
```

```
## [1] 0.00250942
```

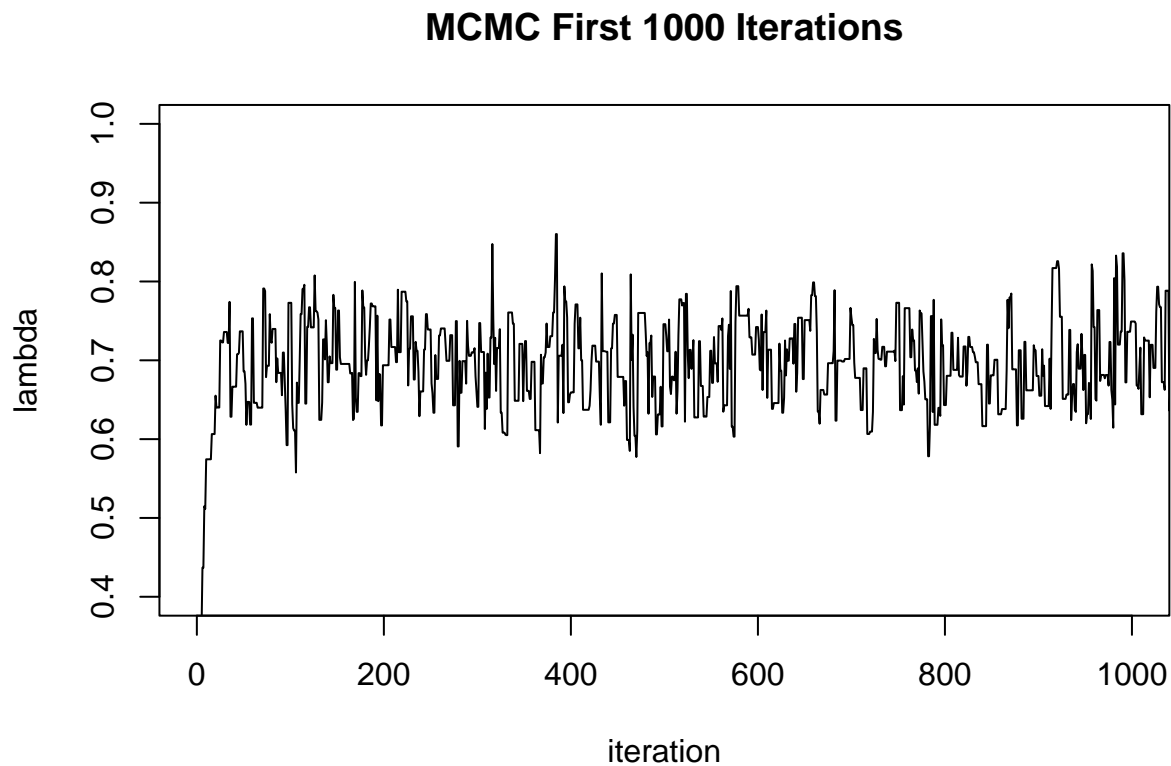
```
print(median)
```

```
##      50%
## 0.7009178
```

```
print(ci)
```

```
##      2.5%      97.5%
## 0.6065487 0.8023146
```

```
plot(chain[,3] , chain[,2] , type = 'l' , ylim = c(0.4 , 1) , xlab = 'iteration' , ylab = 'lambda', xlim
```



Exercise 3

(a) The unbiased frequentist estimator for a Bernoulli process is simply:

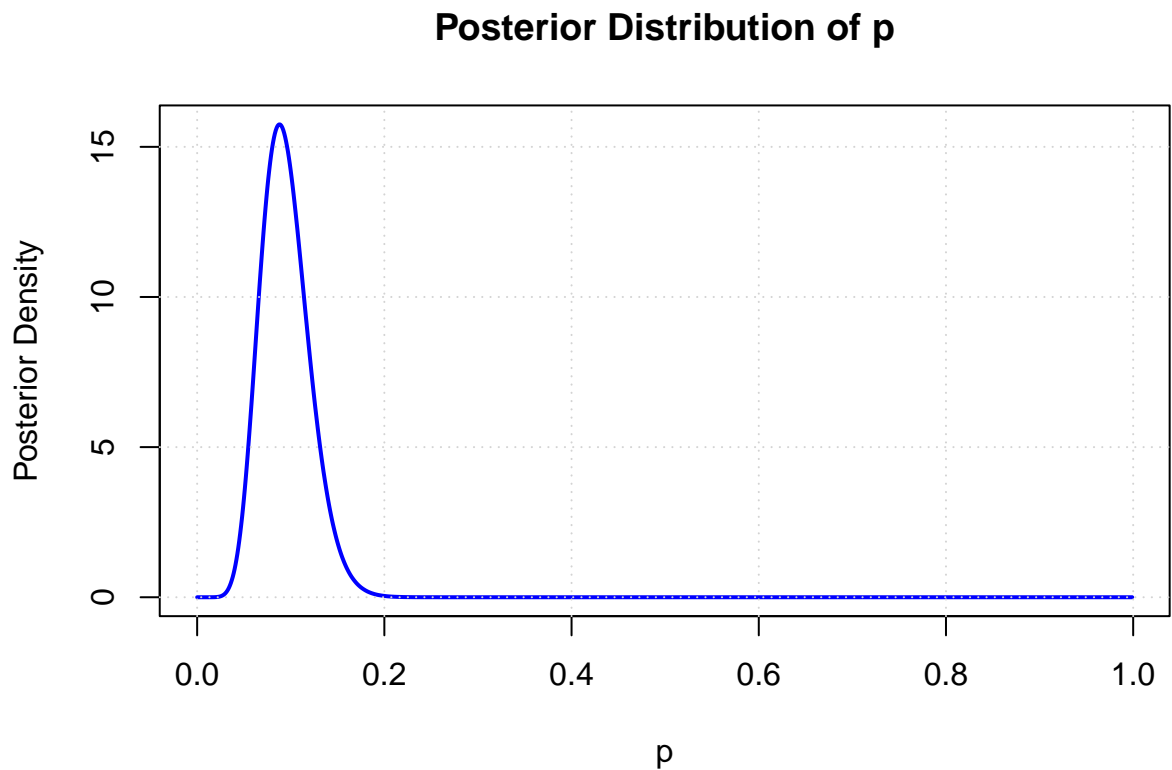
```
y = 11
n = 116
p.hat = y/n
p.hat
```

```
## [1] 0.09482759
```

```
p <- seq(from = 0, by = 0.001, length.out = 1000)
post.p <- dbeta(p, 1 + y, n - y + 10)

# Create the plot
plot(p, post.p, type = 'l', xlab = "p", ylab = "Posterior Density",
     main = "Posterior Distribution of p", col = "blue", lwd = 2)

# Add grid lines
grid()
```



(b)

```
# Adjust the plot margins
par(mar = c(5, 5, 4, 2) + 0.1)
```

```
mean = 0.001 * sum(post.p * p)
variance = 0.001 * sum(p ^ 2 * post.p) - mean ^ 2

sigma = sqrt(variance)

print(mean)
```

(c)

```
## [1] 0.09448819
```

```
print(c(mean - 2 * sigma , mean + 2 * sigma))
```

```
## [1] 0.04277982 0.14619656
```

So the mean of the posterior is considered as the Bayesian estimator for p , and the 95% credibility area is considered as the area 2σ around the mean.

(d) For the frequentist approach, we need to assume the null hypothesis to be true, and using the binomial distribution, find the corresponding p-value :

```
p.value = pbinom(11 , n , 0.1 , lower.tail = T) + pbinom(105 , n , 0.1 , lower.tail = F)
p.value
```

```
## [1] 0.5043125
```

As we see the p-value is much higher than 0.05, so the null hypothesis is not rejected.

For the Bayesian approach, we can use the result of the posterior from the last section. $p = 0.1$ is within the 95% credibility interval, therefore the null hypothesis is not rejected.

(e) We need to to the same task with a new value of y:

```
y = 9
n = 165
p.hat = y/n
p.hat
```

```
## [1] 0.05454545
```


(f) First let's do the Beta prior first:

```
p <- seq(from = 0, by = 0.001, length.out = 1000)
post.p <- dbeta(p, 1 + y, n - y + 10)

# Create the plot
plot(p, post.p, type = 'l', col = 'blue', lwd = 2, xlab = 'p', ylab = 'Posterior Density',
     main = 'Posterior Distribution of p')

# Add grid lines
grid()
```

Posterior Distribution of p



```
# Adjust the plot margins  
par(mar = c(5, 5, 4, 2) + 0.1)
```

```
mean = 0.001 * sum(post.p * p)  
variance = 0.001 * sum(p ^ 2 * post.p) - mean ^ 2  
  
sigma = sqrt(variance)  
  
print(mean)
```

```
## [1] 0.05681818
```

```
print(c(mean - 2 * sigma , mean + 2 * sigma))
```

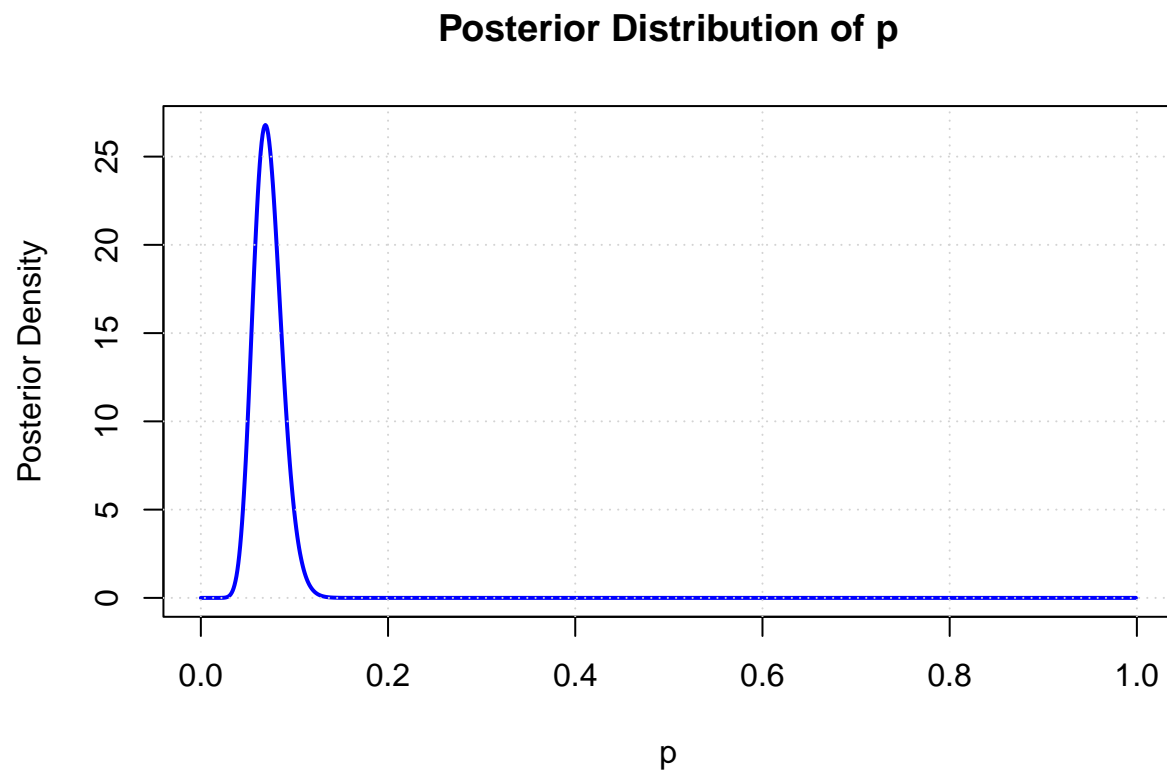
```
## [1] 0.02201774 0.09161862
```

Assuming the previous measurement as a prior we get:

```
p <- seq(from = 0, by = 0.001, length.out = 1000)  
p.prior <- dbeta(p, 1 + 11, 116 - 11 + 10)  
post.p <- dbeta(p, 1 + 11 + 9, 165 - 9 + 116 - 11 + 10)
```

```
# Create the plot
plot(p, post.p, type = 'l', col = 'blue', lwd = 2, xlab = 'p', ylab = 'Posterior Density',
     main = 'Posterior Distribution of p')

# Add grid lines
grid()
```



```
# Adjust the plot margins
par(mar = c(5, 5, 4, 2) + 0.1)
```

```
mean = 0.001 * sum(post.p * p)
variance = 0.001 * sum(p ^ 2 * post.p) - mean ^ 2

sigma = sqrt(variance)

print(mean)
```

(g)

```
## [1] 0.07191781
```

```
print(c(mean - 2 * sigma , mean + 2 * sigma))
```

```
## [1] 0.04173167 0.10210395
```

(h) Frequentist approach:

```
p.value = pbinom(9 , n , 0.1 , lower.tail = T) + pbinom(107 , n , 0.1 , lower.tail = F)
p.value
```

```
## [1] 0.0274604
```

We see that the p-value is smaller than 0.05 and therefore the null hypothesis is rejected.

Bayesian approach: The $p = 0.1$ is almost inside the 95% credibility interval.

Exercise 4

We start by defining the model:

```
library(rjags)
```

```
## Loading required package: coda
```

```
## Linked to JAGS 4.3.1
```

```
## Loaded modules: basemod,bugs
```

```
library(coda)
```

```
data <- list(n = 116, y = 11)
```

```
model <- jags.model("model.txt", data = data)
```

```
## Compiling model graph
```

```
##   Resolving undeclared variables
```

```
##   Allocating nodes
```

```
## Graph information:
```

```
##   Observed stochastic nodes: 1
```

```
##   Unobserved stochastic nodes: 1
```

```
##   Total graph size: 5
```

```
##
```


```
## Initializing model
```

```
samples <- coda.samples(model, variable.names = "p", n.iter = 10000)
```

```
summary(samples)
```

```
##
## Iterations = 1001:11000
## Thinning interval = 1
## Number of chains = 1
## Sample size per chain = 10000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##          Mean          SD      Naive SE Time-series SE
##    0.0938052    0.0253570    0.0002536    0.0003207
##
## 2. Quantiles for each variable:
##
##    2.5%    25%    50%    75%    97.5%
## 0.04961 0.07587 0.09171 0.11009 0.14812
```

```
p.df = as.data.frame(as.mcmc(samples))
hist(p.df$p , breaks = 100 , freq = F , main = 'Histogram of the Samples' , xlab = 'p' )
```



```
mean = mean(p.df$p)
std = sd(p.df$p)
ci = c(mean - 2 * std , mean + 2 * std)

print(mean)
```



```
## [1] 0.09380517
```

```
print(std)
```

```
## [1] 0.02535696
```

```
print(ci)
```

```
## [1] 0.04309125 0.14451908
```